

Individual Assignments

Research Team (3)

Person 1 — LLM Architecture

- Transformers overview
- Attention explained simply

Person 2 — Prompt Interpretation

- How models parse prompts
- Prompt structure analysis

Person 3 — Tokenization + Probabilities

- Tokenization impact
- Next-token prediction basics

Metrics Team (3)

Person 4 — Similarity Metrics

- Cosine similarity
- Embedding distance

Person 5 — Alignment Theory

- Define alignment mathematically
- Confidence scoring ideas

Person 6 — Scoring Designer

- Combine metrics into formula

- Write final scoring proposal

Engineering Team (5)

Person 7 — Repo Architect (Lead)

- Folder structure
- Code standards

Person 8 — Pipeline Builder

- Prompt input handling

Person 9 — Model Interface

- API wrappers or local models

Person 10 — Scoring Integration

- Connect metrics to pipeline

Person 11 — CLI / UI

- Simple CLI or minimal UI

Experiments Team (3)

Person 12 — Dataset Builder

- Prompt categories
- Test sets

Person 13 — Experiment Runner

- Run pipelines
- Log results

Person 14 — Visualization

- Charts
- Graphs
- Insights

Repo Structure

```
/research  
/metrics  
/pipeline  
/experiments  
/docs  
/results
```

Milestones

1. Week 3 — Design Doc
2. Week 5 — First Results
3. Week 8 — Final Release