

Operational Plan - Weekly

Week 1 — Kickoff + Structure

Goal: remove confusion, create clarity

Team lead

- Create repo
- Write README:
 - Vision
 - Tracks
 - Timeline
- Assign tracks
- Pick 4 track leads

Team members

- Read README
- Introduce themselves
- Join channels

Deliverables

- Repo created
 - Track leads confirmed
-

Week 2 — Foundations

Research Track

- LLM architecture summary

- Tokenization basics
- How prompts are processed

Metrics Track

- What does “alignment” mean?
- List 5 scoring ideas
- Read about similarity metrics

Engineering Track

- Repo structure:

```
/research  
/metrics  
/pipeline  
/experiments
```

- Basic Python environment
- Decide libraries

Experiments Track

- Collect:
 - 20 prompts
 - 5 prompt categories

Deliverable: knowledge base exists

Week 3 — Design Week (BIG MILESTONE)

Everything converges here.

Research

- Write final summaries

Metrics

- Propose 2–3 scoring formulas

Engineering

- Define:
 - Input format
 - Output format
 - Pipeline flow

Experiments

- Define evaluation dataset format

Deliverable:

Design Document v1

Week 4 — First Implementation

Engineering

- Build pipeline skeleton
- Load prompts
- Call models (even mock models)

Metrics

- Implement scoring v1

Research

- Support edge cases

Experiments

- Prepare datasets
-

Week 5 — First Results

Your first “wow” moment.

Run:

- Prompt → score pipeline

Output:

- Raw alignment scores
- Known flaws list

This is a huge morale booster.

Week 6 — Iteration Phase

Real research starts here.

Improve:

- Metrics
 - Pipeline speed
 - Better scoring logic
 - Fix inconsistencies
-

Week 7 — Analysis + Insights

Now make it look impressive.

Focus:

- Charts
 - Model comparisons
 - Observations:
 - “Long prompts reduce alignment”
 - “Model X overconfident”
-

Week 8 — Finalization

Deliverables:

- Final repo
- Polished README
- Results folder
- Final presentation